Week 4 Quiz

Bryan Gibson - brg2130

Due Saturday Oct 5 11:59pm

In [1]:
```
import pandas as pd
import seaborn as sns
import statsmodels.api as sm

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)

%matplotlib inline
```

We're going to see if there is a linear relationship between alcohol level and color intensity in wine.

In [2]:
```
# read in ../data/wine_dataset.csv as df
df = pd.read_csv('../data/wine_dataset.csv')
```

In [3]:
```
# using jointplot with kind='reg', plot alcohol (x) against color_int
ensity (y) including a regression line
sns.jointplot(x='alcohol',y='color_intensity',data=df,kind='reg');
```

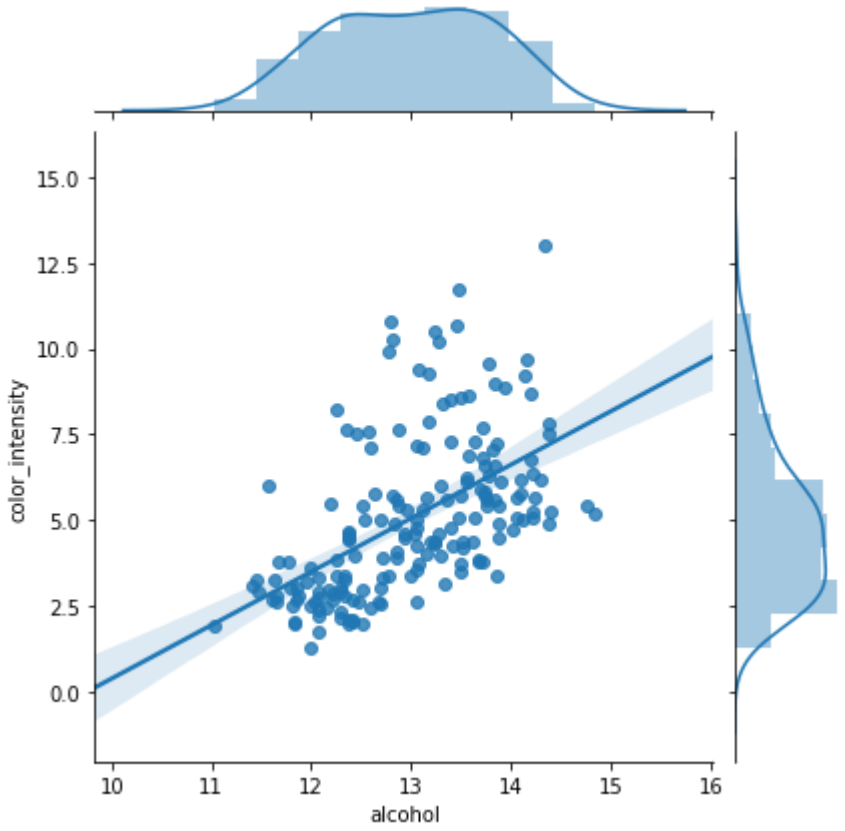

In [4]:
```
# Create a new variable X that is just the alcohol column from df
X = df['alcohol']
```

In [5]:
```
# Add a constant of 1 to X using sm.add_constant (Recall this is our
 intercept or bias term)
X = sm.add_constant(X)
```

In [6]:
```
# Print out the first five rows of X. There should be two columns: co
nst, alcohol
X.iloc[:5]
```

Out[6]:

	const	alcohol
0	1.0	14.23
1	1.0	13.20
2	1.0	13.16
3	1.0	14.37
4	1.0	13.24

In [7]:
```
# Create a new variable y that is just the color_intensity column fro
m df
y = df['color_intensity']
```

In [8]:
```
# Using sm.OLS, instantiate and fit the simple linear model (note: th
e order of paramters is y,X)
model = sm.OLS(y,X).fit()
```

In [9]:
```
# Print out the learned paramters the model
model.params
```

Out[9]:
```
const     -15.225741
alcohol     1.560220
dtype: float64
```

In [10]:
```
# according to the model, what will color_intensity be when alcohol =
0?

# -15.2

# according to the model, does color_intensity go up or down as alcoh
ol increases?

# It goes up.
```

In [11]:
```
# Print out the rsquared value (Recall that this is amount of varianc
e in y explained by the model)
model.rsquared
```

Out[11]:
```
0.2985138336694634
```