

Week 5 Quiz

Perrin Anto - paj2117

```
In [1]: # import the datasets module from sklearn  
from sklearn import datasets
```

```
In [2]: # use datasets.load_boston() to load the Boston housing dataset  
boston = datasets.load_boston()
```

```
In [3]: # print the description of the dataset in boston.DESCR  
print(boston.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.
```

```
:Attribute Information (in order):
```

```
  - CRIM      per capita crime rate by town
  - ZN        proportion of residential land zoned for lots over 2
5,000 sq.ft.
  - INDUS     proportion of non-retail business acres per town
  - CHAS      Charles River dummy variable (= 1 if tract bounds ri
ver; 0 otherwise)
  - NOX       nitric oxides concentration (parts per 10 million)
  - RM        average number of rooms per dwelling
  - AGE       proportion of owner-occupied units built prior to 19
40
  - DIS       weighted distances to five Boston employment centres
  - RAD       index of accessibility to radial highways
  - TAX       full-value property-tax rate per $10,000
  - PTRATIO   pupil-teacher ratio by town
  - B         1000(Bk - 0.63)^2 where Bk is the proportion of black
ks by town
  - LSTAT     % lower status of the population
  - MEDV      Median value of owner-occupied homes in $1000's
```

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

```
This is a copy of UCI ML housing dataset.
```

```
https://archive.ics.uci.edu/ml/machine-learning-databases/housing/
```

```
This dataset was taken from the StatLib library which is maintained at
Carnegie Mellon University.
```

```
The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic
prices and the demand for clean air', J. Environ. Economics & Management,
vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics
...', Wiley, 1980. N.B. Various transformations are used in the table on
pages 244-261 of the latter.
```

```
The Boston house-price data has been used in many machine learning papers
that address regression problems.
```

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

```
In [4]: # copy the dataset features from boston.data to X
X = boston.data
```

```
In [5]: # copy the dataset labels from boston.target to y
y = boston.target
```

```
In [6]: # import the LinearRegression model from sklearn.linear_model
from sklearn.linear_model import LinearRegression
```

```
In [7]: # initialize a linear regression model as lr with the default arguments
lr = LinearRegression()
```

```
In [8]: # fit the lr model using the entire set of X features and y labels
lr.fit(X,y)
```

```
Out[8]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [9]: # score the lr model on entire set of X features and y labels
lr.score(X,y)
```

```
Out[9]: 0.7406426641094095
```

```
In [10]: # import the DecisionTreeRegressor from sklearn.tree
from sklearn.tree import DecisionTreeRegressor
```

```
In [11]: # initialize a decision tree model as dt with the default arguments
dt = DecisionTreeRegressor()
```

```
In [12]: # fit the dt model using the entire set of X features and y labels
dt.fit(X,y)
```

```
Out[12]: DecisionTreeRegressor(criterion='mse', max_depth=None, max_features=None,
                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                presort=False, random_state=None, splitter='best')
```

```
In [13]: # score the dt model on the entire set of X features and y labels  
dt.score(X,y)
```

```
Out[13]: 1.0
```

What are we doing wrong here?!

Why shouldn't we trust these scores to tell us how the models with generalize? We never split the data set into training and testing subsets. These scores are based on the same data that was used to train the model, as seen by the perfect score on the decision tree regressor!